


Markov Decision Process

edge $s \rightarrow a \rightarrow s'$ labelled by $p = P(s' | s, a)$, reward $r = R(s, a, s')$



legend

-  state s
-  hidden state (POMDP)
-  action a
-  observation o
-  transition p, r
-  emission $Z(o | s')$

POMDP — hidden state, observed emission

state is latent; agent sees only o drawn from emission $Z(o | s') = P(o | s')$

