


Multi-head Latent Attention (MLA)

Inference step t — caches only the latent



Key idea: only c_t^{KV} is cached; K and V are reconstructed on the fly, cutting KV-cache memory.

Regular Multi-head Attention (MHA)

Inference step t — caches full K and V

