



Dense feed-forward

one shared FFN — every token



Mixture-of-Experts

router picks top- k of N experts



All N experts share the parameter budget, but only k run per token — so MoE adds capacity at near-constant compute. The router (a learned softmax gate) chooses which.