

The Transformer — encoder–decoder architecture

Vaswani et al., 2017. Garnet = attention sublayers and the encoder memory (K, V) feeding decoder cross-attention.

